

```
import pandas as pd
import numpy as np
```

▼ Lectura de Datos

```
df= pd.read_csv( 'https://raw.githubusercontent.com/PosgradoMNA/Actividades_Aprendizaje-/main/default%20of%20credi

df.head()
```

	ID	X1	X2	X3	X4	X5	X6	X7	X8	X9	...	X15	X16	X17	X18	X19	X20	X21	:
0	1	20000	2.0	2.0	1.0	24.0	2.0	2.0	-1.0	-1.0	...	0.0	0.0	0.0	0.0	689.0	0.0	0.0	
1	2	120000	2.0	2.0	2.0	26.0	-1.0	2.0	0.0	0.0	...	3272.0	3455.0	3261.0	0.0	1000.0	1000.0	1000.0	
2	3	90000	2.0	2.0	2.0	34.0	0.0	0.0	0.0	0.0	...	14331.0	14948.0	15549.0	1518.0	1500.0	1000.0	1000.0	1000.0
3	4	50000	2.0	2.0	1.0	37.0	0.0	0.0	0.0	0.0	...	28314.0	28959.0	29547.0	2000.0	2019.0	1200.0	1100.0	1000.0
4	5	50000	1.0	2.0	1.0	57.0	-1.0	0.0	-1.0	0.0	...	20940.0	19146.0	19131.0	2000.0	36681.0	10000.0	9000.0	689.0

5 rows x 25 columns

▼ Obten la información del DataFrame con los métodos y propiedades: shape, columns, head(), dtypes, info(), isna().

```
df.shape
```

```
(30000, 25)
```

```
df.columns
```

```
Index(['ID', 'X1', 'X2', 'X3', 'X4', 'X5', 'X6', 'X7', 'X8', 'X9', 'X10',
      'X11', 'X12', 'X13', 'X14', 'X15', 'X16', 'X17', 'X18', 'X19', 'X20',
      'X21', 'X22', 'X23', 'Y'],
      dtype='object')
```

```
df.dtypes
```

```
ID      int64
X1      int64
X2     float64
X3     float64
X4     float64
X5     float64
X6     float64
X7     float64
X8     float64
X9     float64
X10    float64
X11    float64
X12    float64
X13    float64
X14    float64
X15    float64
X16    float64
X17    float64
X18    float64
X19    float64
X20    float64
X21    float64
X22    float64
X23    float64
Y      float64
dtype: object
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 30000 entries, 0 to 29999
Data columns (total 25 columns):
#   Column  Non-Null Count  Dtype
---  -
0    ID      30000 non-null    int64
1    X1      30000 non-null    int64
2    X2      29999 non-null    float64
```

```
3  X3      29998 non-null float64
4  X4      29998 non-null float64
5  X5      29995 non-null float64
6  X6      29997 non-null float64
7  X7      29995 non-null float64
8  X8      29993 non-null float64
9  X9      29991 non-null float64
10 X10     29984 non-null float64
11 X11     29986 non-null float64
12 X12     29989 non-null float64
13 X13     29989 non-null float64
14 X14     29987 non-null float64
15 X15     29985 non-null float64
16 X16     29983 non-null float64
17 X17     29990 non-null float64
18 X18     29992 non-null float64
19 X19     29991 non-null float64
20 X20     29992 non-null float64
21 X21     29989 non-null float64
22 X22     29989 non-null float64
23 X23     29995 non-null float64
24 Y       29997 non-null float64
dtypes: float64(23), int64(2)
memory usage: 5.7 MB
```

df.isna()

	ID	X1	X2	X3	X4	X5	X6	X7	X8	X9	...	X15	X16	X17	X18	X19	X20	X21
0	False	False	False	False	False	False	False	False	False	False	...	False	False	False	False	False	False	False
1	False	False	False	False	False	False	False	False	False	False	...	False	False	False	False	False	False	False
2	False	False	False	False	False	False	False	False	False	False	...	False	False	False	False	False	False	False
3	False	False	False	False	False	False	False	False	False	False	...	False	False	False	False	False	False	False
4	False	False	False	False	False	False	False	False	False	False	...	False	False	False	False	False	False	False
...
29995	False	False	False	False	False	False	False	False	False	False	...	False	False	False	False	False	False	False
29996	False	False	False	False	False	False	False	False	False	False	...	False	False	False	False	False	False	False
29997	False	False	False	False	False	False	False	False	False	False	...	False	False	False	False	False	False	False
29998	False	False	False	False	False	False	False	False	False	False	...	False	False	False	False	False	False	False
29999	False	False	False	False	False	False	False	False	False	False	...	False	False	False	False	False	False	False

30000 rows x 25 columns

df.isnull().values.any()

True

df.isnull().any()

```
ID      False
X1      False
X2       True
X3       True
X4       True
X5       True
X6       True
X7       True
X8       True
X9       True
X10      True
X11      True
X12      True
X13      True
X14      True
X15      True
X16      True
X17      True
X18      True
X19      True
X20      True
X21      True
X22      True
X23      True
Y        True
dtype: bool
```

df.isna().any()

```
ID      False
X1      False
X2       True
X3       True
X4       True
X5       True
X6       True
X7       True
X8       True
X9       True
X10      True
X11      True
X12      True
X13      True
X14      True
X15      True
X16      True
X17      True
X18      True
X19      True
X20      True
X21      True
X22      True
X23      True
Y        True
dtype: bool
```

▼ Limpia los datos eliminando los registros nulos o rellena con la media de la columna.

```
df.dropna(inplace = True)
df
```

	ID	X1	X2	X3	X4	X5	X6	X7	X8	X9	...	X15	X16	X17	X18	X19	X20	
0	1	20000	2.0	2.0	1.0	24.0	2.0	2.0	-1.0	-1.0	...	0.0	0.0	0.0	0.0	689.0	0.0	
1	2	120000	2.0	2.0	2.0	26.0	-1.0	2.0	0.0	0.0	...	3272.0	3455.0	3261.0	0.0	1000.0	1000.0	1000.0
2	3	90000	2.0	2.0	2.0	34.0	0.0	0.0	0.0	0.0	...	14331.0	14948.0	15549.0	1518.0	1500.0	1000.0	1000.0
3	4	50000	2.0	2.0	1.0	37.0	0.0	0.0	0.0	0.0	...	28314.0	28959.0	29547.0	2000.0	2019.0	1200.0	1000.0
4	5	50000	1.0	2.0	1.0	57.0	-1.0	0.0	-1.0	0.0	...	20940.0	19146.0	19131.0	2000.0	36681.0	10000.0	9000.0
...
29995	29996	220000	1.0	3.0	1.0	39.0	0.0	0.0	0.0	0.0	...	88004.0	31237.0	15980.0	8500.0	20000.0	5003.0	3000.0
29996	29997	150000	1.0	3.0	2.0	43.0	-1.0	-1.0	-1.0	-1.0	...	8979.0	5190.0	0.0	1837.0	3526.0	8998.0	0.0
29997	29998	30000	1.0	2.0	2.0	37.0	4.0	3.0	2.0	-1.0	...	20878.0	20582.0	19357.0	0.0	0.0	22000.0	4000.0
29998	29999	80000	1.0	3.0	1.0	41.0	1.0	-1.0	0.0	0.0	...	52774.0	11855.0	48944.0	85900.0	3409.0	1178.0	1000.0
29999	30000	50000	1.0	2.0	1.0	46.0	0.0	0.0	0.0	0.0	...	36535.0	32428.0	15313.0	2078.0	1800.0	1430.0	1000.0

29975 rows x 25 columns

▼ Calcula la estadística descriptiva con describe() y explica las medidas de tendencia central y dispersión.

```
df.describe()
```

	ID	x1	x2	x3	x4	x5	x6	x7	
count	29975.000000	29975.000000	29975.000000	29975.000000	29975.000000	29975.000000	29975.000000	29975.000000	2
mean	14999.137581	167538.604837	1.604070	1.853078	1.551827	35.485486	-0.016981	-0.134012	
std	8657.789302	129742.083982	0.489058	0.790468	0.521994	9.217725	1.123791	1.197162	

df.shape

(29975, 25)

50%

14999 0000000

140000 0000000

2 0000000

2 0000000

2 0000000

34 0000000

0 0000000

0 0000000

▼ Realiza el conteo de las variables categóricas.

```
numericas=df.drop(['ID','x2','x3','x4','x6','x7','x8','x9','x10','x11','Y'], axis=1)
```



```
numericas.shape
```

(29975, 14)

El conteo de variables categoricas es el numero de variables total menos el numero de variables categoricas. El dataframe numericas solo queda conformado por 14 variables numericas. Por lo tanto el conteo de variables categoricas es 11.

▼ Escala los datos, si consideras necesario.

```
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
scaled = scaler.fit_transform(numericas)
#Son muchos valores. Así que imprimamos los primeros 5 resultados mejor.
scaled[:5]
```

```
array([[ -1.13718742, -1.24604254, -0.64248634, -0.6473368 , -0.66789986,
        -0.67243242, -0.66304357, -0.65268019, -0.34190246, -0.22712701,
        -0.29676957, -0.30811617, -0.31414687, -0.29346219],
       [-0.36641463, -1.02906565, -0.6591989 , -0.6666781 , -0.63916972,
        -0.621586  , -0.6062293 , -0.59793705, -0.34190246, -0.21363445,
        -0.23999307, -0.24430798, -0.31414687, -0.1810025 ],
       [-0.59764647, -0.1611581 , -0.29865028, -0.49388461, -0.48237208,
        -0.4497307 , -0.41723755, -0.39165562, -0.25028162, -0.19194222,
        -0.23999307, -0.24430798, -0.24871751, -0.01231298],
       [-0.90595558,  0.16430723, -0.0576555 , -0.01342823,  0.03272337,
        -0.23243685, -0.18683962, -0.15666804, -0.22118989, -0.16942569,
        -0.22863777, -0.23792716, -0.24420288, -0.23723234],
       [-0.90595558,  2.33407612, -0.57862293, -0.61126676, -0.1612519 ,
        -0.34702777, -0.34820532, -0.33152378, -0.22118989,  1.334366  ,
         0.27099537,  0.2661576 , -0.26906604, -0.25528212]])
```

```
scaled_df = pd.DataFrame(scaled, columns=numericas.columns)
scaled_df.head()
```

	x1	x5	x12	x13	x14	x15	x16	x17	x18	x19	x20	
0	-1.137187	-1.246043	-0.642486	-0.647337	-0.667900	-0.672432	-0.663044	-0.652680	-0.341902	-0.227127	-0.296770	-0.308116
1	-0.366415	-1.029066	-0.659199	-0.666678	-0.639170	-0.621586	-0.606229	-0.597937	-0.341902	-0.213634	-0.239993	-0.293462
2	-0.597646	-0.161158	-0.298650	-0.493885	-0.482372	-0.449731	-0.417238	-0.391656	-0.250282	-0.191942	-0.239993	-0.248718
3	-0.905956	0.164307	-0.057655	-0.013428	0.032723	-0.232437	-0.186840	-0.156668	-0.221190	-0.169426	-0.228638	-0.237927
4	-0.905956	2.334076	-0.578623	-0.611267	-0.161252	-0.347028	-0.348205	-0.331524	-0.221190	1.334366	0.270995	0.266158

▼ Reduce las dimensiones con PCA, si consideras necesario.

```
from sklearn.decomposition import PCA
pcs = PCA()
pcs_t = pcs.fit_transform(scaled_df)

pcsSummary_df = pd.DataFrame({'% varianza explicada': np.round(pcs.explained_variance_ratio_,4) * 100,
'% varianza acumulada': np.cumsum(pcs.explained_variance_ratio_) * 100})
pcsSummary_df
```

	% varianza explicada	% varianza acumulada
0	42.28	42.279619
1	12.26	54.537070
2	7.46	62.001690
3	6.61	68.608539
4	6.31	74.920341
5	6.23	81.153789
6	5.57	86.722019
7	5.19	91.912021
8	5.05	96.963223
9	1.89	98.852053
10	0.51	99.359387
11	0.29	99.652813
12	0.18	99.833971
13	0.17	100.000000

```
pcs_labels = [f'PC{i + 1}' for i in range(len(scaled_df.columns))]  
pcsSummary_df.index = pcs_labels  
pcsSummary_df
```

	% varianza explicada	% varianza acumulada
PC1	42.28	42.279619
PC2	12.26	54.537070
PC3	7.46	62.001690
PC4	6.61	68.608539
PC5	6.31	74.920341
PC6	6.23	81.153789
PC7	5.57	86.722019
PC8	5.19	91.912021
PC9	5.05	96.963223
PC10	1.89	98.852053
PC11	0.51	99.359387
PC12	0.29	99.652813
PC13	0.18	99.833971
PC14	0.17	100.000000

```
pcs_df = pd.DataFrame(pcs_t, columns =pcs_labels)  
print("Varianza total variables originales: ", scaled_df.var().sum())  
print("Varianza total de los componentes: ", pcs_df.var().sum())  
  
pcsSummary_df
```

Varianza total variables originales: 14.000467071461934
Varianza total de los componentes: 14.000467071461943

	% varianza explicada	% varianza acumulada
PC1	42.28	42.279619
PC2	12.26	54.537070

- Indica la varianza de los datos explicada por cada componente seleccionado.
- Para actividades de exploración de los datos la varianza > 70%

```
total_var =scaled_df.var().sum()
pd.DataFrame({"Porcentaje Varianza": (scaled_df.var()/ total_var) * 100,"Porcentaje Varianza Acumulado": (scaled_d
```

	Porcentaje Varianza	Porcentaje Varianza Acumulado
X1	7.142857	7.142857
X5	7.142857	14.285714
X12	7.142857	21.428571
X13	7.142857	28.571429
X14	7.142857	35.714286
X15	7.142857	42.857143
X16	7.142857	50.000000
X17	7.142857	57.142857
X18	7.142857	64.285714
X19	7.142857	71.428571
X20	7.142857	78.571429
X21	7.142857	85.714286
X22	7.142857	92.857143
X23	7.142857	100.000000

- Indica la importancia de las variables en cada componente

```
comps_df = pd.DataFrame(
pcs.components_.round(4),
columns = pcs_df.columns,
index = scaled_df.columns)
comps_df.iloc[:, :7]

comps_df.iloc[:, :7].abs().idxmax()
comps_df.iloc[:, :10].abs().idxmax()

PC1      X18
PC2      X12
PC3      X19
PC4      X23
PC5      X23
PC6      X20
PC7      X22
PC8      X20
PC9      X17
PC10     X17
dtype: object


comps_df
```

	PC1	PC2	PC3	PC4	PC5	PC6	PC7	PC8	PC9	PC10	PC11	PC12	PC13	PC14
X1	0.1655	0.0327	0.3724	0.3832	0.3883	0.3915	0.3885	0.3807	0.1352	0.1168	0.1281	0.1169	0.1138	0.1055
X5	0.3008	0.0718	-0.1909	-0.1747	-0.1269	-0.1204	-0.1060	-0.0942	0.3833	0.4083	0.3923	0.3495	0.3041	0.3235
X12	-0.3785	-0.8697	-0.0342	-0.0018	0.0347	0.0340	0.0339	0.0185	0.1735	0.2008	0.1220	0.0622	-0.0605	-0.0507
X13	-0.2004	0.3386	-0.0640	0.0073	0.0605	0.0748	0.0397	-0.0703	0.3613	0.3464	0.2453	-0.0942	-0.6089	-0.3672
X14	0.0347	-0.0390	0.0411	0.0831	0.1142	0.0286	-0.1070	-0.1649	0.2262	0.1506	-0.2392	-0.5795	-0.1928	0.6573
X15	-0.0783	0.0711	-0.0440	-0.0290	0.0988	0.0144	-0.0990	0.0698	0.0398	0.4072	-0.1079	-0.4992	0.6036	-0.4110
X16	0.1112	-0.0787	0.0082	-0.0323	-0.1213	0.1264	-0.0076	0.0080	-0.2011	-0.2796	0.7852	-0.4621	0.0147	0.0253
X17	-0.0493	0.0285	0.0095	-0.1357	0.0928	0.0392	0.0498	0.0001	-0.7490	0.5778	0.0695	0.0777	-0.1631	0.1825


Elabora los histogramas de los atributos para visualizar su distribución

```
X20 -0.0062  0.0001  0.4159  0.0384 -0.4845 -0.5233  0.0680  0.5136  0.0476  0.1472  0.0002 -0.1158 -0.0995  0.0350

hist = comps_df.hist(bins=6)
```



```
boxplot = comps_df.boxplot(column=['PC1', 'PC2', 'PC3', 'PC4','PC5', 'PC6', 'PC7', 'PC8','PC9', 'PC10'])
```



```
from matplotlib import pyplot as plt

Language = ['PC1', 'PC2', 'PC3', 'PC4','PC5', 'PC6', 'PC7', 'PC8','PC9', 'PC10']
data = [0.1655, 0.0328, 0.3724, 0.3833, 0.3883, 0.3916, 0.3885, 0.3807, 0.1351, 0.1168]
# Creating plot
fig = plt.figure(figsize =(10, 7))
plt.pie(data, labels = Language)
# show plot
plt.show()
```





Interpreta y explica cada uno de los gráficos indicando cuál es la información más relevante que podría ayudar en el proceso de toma de decisiones.



Rich text editor toolbar with icons for bold, italic, code, link, image, list, and other formatting options.

Como se puede observar en el histograma, las variables más representativas después de aplicar la técnica de PCA tienen una menor varianza y su distribución es uniforme. Lo mismo se puede observar con la gráfica de Boxplot. En el último gráfico de pie, podemos ver cómo están distribuidos los pesos de cada componente que se analizó. Con ello se podemos eliminar casi la mitad de variables del set de datos original.

Como se puede observar en el histograma, las variables más representativas después de aplicar la técnica de PCA tienen una menor varianza y su distribución es uniforme. Lo mismo se puede observar con la gráfica de Boxplot. En el último gráfico de pie, podemos ver cómo están distribuidos los pesos de cada componente que se analizó. Con ello se podemos eliminar casi la mitad de variables del set de datos original.